

## APPENDIX B

## TLA+ FORMAL SPECIFICATION AND VERIFICATION

To rigorously validate the safety and liveness properties of AdaptiveBFT, particularly its *Adaptive View-Change* (AVC) and *Adaptive Pipeline Scheduling* (APS) mechanisms, we developed a comprehensive TLA+ formal specification and checked it with TLC.

## A. AdaptiveBFT Formalization Structure

A formal specification of AdaptiveBFT appears in Figure 10, with supporting definitions in Figures 11 to 13 and correctness properties in Figure 14.

The split follows a five-figure, reviewer-friendly structure:

- 1) Main protocol state-machine excerpt (Fig. 10).
- 2) Message/proposal/state type system (Fig. 11).
- 3) AVC + RVS support definitions (Fig. 12).
- 4) APS + mempool support definitions (Fig. 13).
- 5) Correctness properties only (Fig. 14).

In addition to this split view, we also provide two complementary rendered artifacts in the repository root/module folder: `AdaptiveBFT.pdf` (monolithic consolidated rendering) and `APS_Scheduler.pdf` (scheduler-focused support rendering).

## B. Modeling Scope and Abstractions

To keep exhaustive TLC exploration tractable without obscuring the core protocol logic, we make the following explicit abstractions:

- A1: Bounded BFT domain.** The baseline model instantiates  $f = 1$ ,  $n = 3f + 1 = 4$ , and  $q = 2f + 1 = 3$ , with finite bounds on views, batching, pipeline depth, timeout, reputation range, and transaction set.
- A2: Typed message/proposal semantics.** Protocol artifacts (`MINORMESSAGE`, `FULLMESSAGE`, `TEPROPOSAL`, `VPROPOSAL`, `NPMESSAGE`, `SYNMESSAGE`) are represented as typed TLA+ records.
- A3: Cryptographic abstraction.** QC checks and RVS/VRF behavior are encoded as deterministic typed operators over bounded domains with explicit witness validation; cryptographic hardness assumptions remain out of model scope.
- A4: Timer and network abstraction.** Timeout and delay effects are represented via symbolic state (`stageTimer`, `timeout votes`, `Stable/Unstable`) rather than packet-level or wall-clock models.
- A5: Adaptive adversary abstraction.** Adaptive corruption is modeled as bounded perturbation transitions parameterized by `MaxAttackCount` in attack-enabled configurations.
- A6: Reputation-game abstraction.** Bayesian-report normalization, finite PageRank-style updates, and payoff/cost-based truthful-report utilities are encoded as finite operators over bounded report/reputation domains.
- A7: APS control-plane abstraction.** Adaptive scheduling is modeled as a finite reconfiguration loop over five abstract states: `Monitor`, `Sample`, `Estimate`, `Explore`, and `Deploy`, with finite candidate configurations.

**A8: Eventual synchrony abstraction.** Liveness is checked under explicit fairness constraints in the liveness wrappers, capturing eventual progress after unstable episodes.

**A9: Safety-wrapper quorum compression.** In `MC_AdaptiveAttack`, per-replica vote/ack interleavings are collapsed into quorum-completion steps to keep adaptive-attack safety exploration tractable.

These assumptions define the exact claim boundary of the model-checking results reported below.

## C. Checked Properties and TLC Results

The correctness profile in Figure 14 consists of 13 safety invariants (I1–I13), an aggregate Safety formula, and temporal liveness goals (Liveness and related progress formulas).

*a) Baseline claim (`MaxView=2`):* Using the TLC model-checker, we are able to check that all possible executions of the bounded AdaptiveBFT model in a 4-node system with one Byzantine node satisfy the safety and liveness properties in Figure 14 under the eventual-synchrony abstraction (encoded by fairness constraints), for a maximum of 3 protocol rounds (views 0..2). In this baseline setting, the full campaign (`MC_AdaptiveAttack`, `MC_LivenessAPS`, `MC_LivenessAPSAttack`) completes in about 26 seconds on an AMD Ryzen 7 7840H platform (16 vCPUs, 15 GiB RAM, Java -Xmx8G), with no TLC-reported invariant or temporal-property violations.

*b) Strengthened claim (`MaxView=4`):* To increase depth coverage, we provide a strengthened profile that changes only `MaxView` from 2 to 4 (all other constants fixed). Under this profile, TLC reports no invariant or temporal-property violation: `MC_AdaptiveAttack` explores 113,041 generated states (32,774 distinct) in 21s, `MC_LivenessAPS` explores 16,791 generated states (5,047 distinct) in 39s, and `MC_LivenessAPSAttack` explores 104,191 generated states (33,269 distinct) in 55s.

*c) Refinement-transfer diagnostics:* To reduce the remaining gap between bounded concrete exploration and the abstract theorem layer, we additionally check a dedicated transfer wrapper (`MC_RefinementTransfer`) that enforces `StepProjectionChecked` and `StepBoxProjectionChecked` under a tractable state-space constraint. In baseline/`MaxView=3`/`MaxView=4` profiles, TLC reports no transfer-property violation and explores 269,339 generated states (47,796 distinct, diameter 30), with runtime 1m10s/1m21s/1m21s, respectively. In addition, we embed wrapper-level transfer checks directly in `MC_AdaptiveAttack`, `MC_LivenessAPS`, and `MC_LivenessAPSAttack`; these also pass in baseline and `MaxView=4` runs.

Detailed state-space scale and per-model runtime are reported in Tables II to X.

**Local variables:**transition families  $M_1, M_2, M_3, M_4$ : $M_1$ : ingress + pre-validation $M_2$ : APS control loop $M_3$ : consensus pipeline $M_4$ : AVC + RVS + sync update $Next \triangleq M_1 \vee M_2 \vee M_3 \vee M_4$ 

```

1: procedure MOne                                ▷ MainFlow:9–12
2:   return InjectTxStep ∨ PreValidateStep ∨ AgeTxStep
3: procedure MTwo                                ▷ MainFlow:14–19
4:   return DetectAnomaly ∨ SampleGrid ∨ EstimateGrid
   ∨ ExploreGrid ∨ DeployConfig

5: procedure MThree                               ▷ MainFlow:21–30
6:   return PreOrder ∨ GenTentative ∨ RecoverFull
   ∨ PrepareVote ∨ PrepareQC
   ∨ PreCommitVote ∨ PreCommitQC
   ∨ CommitVote ∨ DecideBlock

7: procedure MFour                                ▷ MainFlow:32–40
8:   return Tick ∨ CastTimeoutVote ∨ StartViewChange
   ∨ BroadcastNP ∨ ConfirmNewPrimary
   ∨ BroadcastSyn ∨ SendSynAck
   ∨ CompleteViewChange

9: procedure Next                                ▷ MainFlow:42–46
10:  return MOne ∨ MTwo ∨ MThree ∨ MFour

```

Fig. 10: High-level TLA+ state-machine excerpt of AdaptiveBFT.

**Local variables:**

universes:  $MsgType, ConsensusPhase, SchedulerStateType, NetworkConditionType$   
 record tags: Minor, Full, TeProposal, ReProposal, VProposal, NPMesssage, SynMessage  
 helper operators:  $QC(view), NilQC, SamePrefix$

```

1: procedure MkMinor( $v, a, q, s$ )                ▷ Types:23–24
2:   return ⟨Minor,  $v, a, q, sprocedure MkFull( $v, tx, q, p, s$ )              ▷ Types:26–27
4:   return ⟨Full,  $v, tx, q, p, sprocedure MkTe( $v, al, q, p, s$ )
6:   return ⟨TeProposal,  $v, al, q, p, sprocedure MkRe( $v, tx, q, p, s$ )
8:   return ⟨ReProposal,  $v, tx, q, p, sprocedure MkV( $v, rv, q, p, s$ )
10:  return ⟨VProposal,  $v, rv, q, p, sprocedure MkNP( $v, l, t, k, \pi, q, s$ )
12:  return ⟨NPMesssage,  $v, l, t, k, \pi, q, sprocedure MkSyn( $v, l, rv, q, s$ )
14:  return ⟨SynMessage,  $v, l, rv, q, sprocedure SamePrefix( $s_1, s_2$ )
16:   $m \leftarrow MinNat(Len(s_1), Len(s_2))$ 
17:  return  $\forall i \in 1..m : s_1[i] = s_2[i]$$$$$$$$ 
```

Fig. 11: TLA+ type and record definitions for messages, proposals, and protocol state.

**Local variables:**

$rep[n]$ : bounded reputation  
 $thr, view, \kappa$ : threshold, epoch, strike parameter  
 evidence tuple  $ev = \langle ticket, strikes, proof \rangle$

```

1: procedure CandSet( $rep, thr$ )                  ▷ AVC_RVS:38–46
2:    $base \leftarrow \{n : rep[n] \geq thr\}$ 
3:    $hi \leftarrow \{n \in base : rep[n] \geq AverageRep(rep)\}$ 
4:   return  $hi$  if non-empty; else  $base$  if non-empty; else all nodes

5: procedure RVSCContext( $rep, thr$ )               ▷ AVC_RVS:56–65
6:    $cand \leftarrow CandSet(rep, thr)$ 
7:   return ⟨cand, SumRep( $rep, cand$ ), RVSKappa( $cand, thr$ ), max( $rep$ )⟩

8: procedure SelectLeader( $rep, thr, view$ )          ▷ AVC_RVS:67–86
9:    $ctx \leftarrow RVSCContext(rep, thr)$ 
10:   $wins \leftarrow ValidWinners(rep, ctx, view)$ 
11:  return PickWinner( $rep, wins, ctx, view$ )

12: procedure LeaderEvidence( $rep, thr, view, ldr$ ) ▷ AVC_RVS:88–100
13:    $ctx \leftarrow RVSCContext(rep, thr)$ 
14:    $r \leftarrow SortitionResult(rep, ldr, ctx, view)$ 
15:   return ⟨r.ticket, r.strikes, r.proof⟩

16: procedure VerifyLeader( $rep, thr, view, ldr, ev$ ) ▷ AVC_RVS:102–139
17:    $ctx \leftarrow RVSCContext(rep, thr)$ 
18:    $r \leftarrow SortitionResult(rep, ldr, ctx, view)$ 
19:    $ok \leftarrow SortitionVerify(rep, ldr, ctx, view, r)$ 
20:   return  $(ldr \in ctx.cand) \wedge ok$ 
    $\wedge (ev = \langle r.ticket, r.strikes, r.proof \rangle)$ 
21: procedure DecayUpdate( $rep, n, h, m, a, b$ )      ▷ AVC_RVS:141–160
22:    $obs \leftarrow m$  if  $h$  else 0
23:   return DecayUpdateByObservation(...)  

   with arguments  $(rep, n, obs, m, a, b)$ 

```

Fig. 12: TLA+ support definitions for AVC and RVS.

**Local variables:**

```

APS config  $cfg = [batchSize, pipelineDepth, timeout]$ 
scheduler state  $\in \{\text{Monitor}, \text{Sample}, \text{Estimate}, \text{Explore}, \text{Deploy}\}$ 
mempool state:  $pool, age, hotTx, warmTx, agingThreshold$ 

1: procedure NetFeasible( $cfg, net$ ) ▷ APS_Mempool:17–24
2:   if  $net = \text{Unstable}$  then
3:     return  $(cfg.timeout \geq 2) \wedge (cfg.batchSize \leq 2) \wedge (cfg.pipelineDepth \leq 2)$ 
4:   else
5:     return  $(cfg.timeout \geq 1) \wedge (cfg.batchSize \geq 1) \wedge (cfg.pipelineDepth \geq 1)$ 
6: procedure Perf( $cfg, net$ ) ▷ APS_Mempool:26–40
7:   return  $2 \cdot \text{LatencyScore}(cfg, net) - \text{ThroughputScore}(cfg)$ 
8: procedure NextSched( $st$ ) ▷ APS_Mempool:47–52
9:   return Monitor  $\rightarrow$  Sample  $\rightarrow$  Estimate  $\rightarrow$  Explore  $\rightarrow$  Deploy  $\rightarrow$  Monitor
10: procedure Prio( $tx, age, hot, warm, th$ ) ▷ APS_Mempool:56–59
11:   return High if  $(tx \in hot) \vee (age[tx] \geq th)$ 
     else Mid if  $(tx \in warm)$  else Low
12: procedure Front( $pool, age, hot, warm, th$ ) ▷ APS_Mempool:65–93
13:    $(hi, mid, lo) \leftarrow \text{partition by } \text{Prio}(\cdot)$ 
14:   return  $hi$  if non-empty; else  $mid$  if non-empty; else  $lo$ 
15: procedure PickBatch( $pool, age, hot, warm, th, lim$ ) ▷ APS_Mempool:95–100
16:    $front \leftarrow \text{Front}(pool, age, hot, warm, th)$ 
17:   return non-empty  $chosen \subseteq front$  with  $|chosen| \leq lim$ 
18: procedure Recover( $abst, valid$ ) ▷ APS_Mempool:102–103
19:   return  $abst \cap valid$ 
20: procedure BumpAge( $age, tx, maxAge$ ) ▷ APS_Mempool:105–109
21:   return bounded age increment; reset on commit

```

Fig. 13: TLA+ support definitions for APS and mempool processing.

**Local variables:**

```

safety invariants:  $I1, \dots, I13$ 
liveness goals:  $L1$  (one commit),  $L2$  (two commits),  $L3$  (infinite collect),  $L4$  (view progress)
macros: Safety, Liveness, wrapper checks

1: procedure SafetyProfile ▷ Properties:22–35
2:   return  $(I1 \wedge I2 \wedge I3 \wedge I4 \wedge I5 \wedge I6)$ 
      $\wedge (I7 \wedge I8 \wedge I9 \wedge I10 \wedge I11 \wedge I12 \wedge I13)$ 
3: procedure Safety ▷ Properties:44
4:   return SafetyProfile

5: procedure Liveness ▷ Properties:48
6:   return  $L2$ 
7: procedure LiveSet ▷ Properties:37–40
8:   return  $\{L1, L2, L3, L4\}$ 

9: procedure ClaimBoundary ▷ explicit boundary
10:  return finite-domain TLC checks
    + fairness-based eventual synchrony
    + bounded adaptive perturbation model
11: procedure ReportedClaim
12:  return all explored executions satisfy safety and liveness

```

Fig. 14: TLA+ formalization of AdaptiveBFT safety and liveness properties.

TABLE II: AdaptiveBFT verification settings and checked property profile.

Model	Assumptions	Invariants (abbr.)	Liveness target
MC_AdaptiveAttack	$f = 1, n = 4, q = 3$ ; Node = $\{r1, r2, r3, r4\}$ ; bounded transaction/pipeline/view/timeout/reputation domains; MaxAttackCount=1; quorum-compressed vote/ack transitions	I1–I13 + B1–B3 <sup>a,d</sup>	W-A <sup>*c</sup>
MC_LivenessAPS	Same bounded domain as left (without attack bound); fairness-structured liveness wrapper	I1–I13 + B1–B3 <sup>a,d</sup>	L1 + W-P <sup>*c</sup>
MC_LivenessAPSAttack	Same bounded domain as left + MaxAttackCount=1; joint fairness and recovery wrapper	I1–I13 + B1–B3 <sup>a,d</sup>	L2 + W-J <sup>*c</sup>
MC_RefinementTransfer	Same bounded constants as baseline/mv3/mv4; transfer constraint (schedulerState=Monitor, stable network, initial view/chain)	I1–I13 + B1–B3 <sup>a,d</sup>	T0–T4 <sup>b</sup>

<sup>a</sup> I1 TypeOK; I2 Consistency; I3 NoForkPerView; I4 PipelineBounded; I5 QCLocked; I6 QCViewSafety; I7 MempoolSoundness; I8 ProposalFlowSafety; I9 ViewChangeMessageSafety; I10 ChainParentSafety; I11 PrimaryEligibilitySafety; I12 ReconfigurationSafety; I13 DecoupledPipelineSafety.

<sup>b</sup> Transfer checks: T0 InitProjectionChecked; T1 StepProjectionChecked; T2 StepBoxProjectionChecked; T3 SpecProjectionChecked; T4 SpecToAbstractSpec-Checked.

<sup>c</sup> Wrapper projection families: W-A\* = attack wrapper projection checks (Init/Step/StepBox/Spec/SpecToAbstractSpec); W-P\* and W-J\* denote the same five check kinds for APS and Joint wrappers.

<sup>d</sup> Wrapper finite bridge diagnostics: B1 RefinementInvariantCore\*; B2 SafetyBridgeFinite\*; B3 CommitProjectionShape\* (Attack/APS/Joint/Transfer suffix by model).

TABLE III: AdaptiveBFT deepened-profile TLC outcomes (MaxView=5, bounded full-suite).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack	299,733	86,886	80	01min 42s	Pass	Safety + W-A*
MC_LivenessAPS	20,143	6,053	115	01min 14s	Pass	Safety + L1 + W-P*
MC_LivenessAPSAttack	143,795	45,553	124	01min 39s	Pass	Safety + L2 + W-J*
MC_RefinementTransfer	269,339	47,796	30	01min 27s	Pass	T0–T4

Snapshot source: make suite-mv5 (artifact: docs/results/mv5\_suite\_latest.md, executed on 2026-02-21). This is bounded finite-domain evidence with deeper view depth; it is not an unbounded theorem proof.

TABLE IV: AdaptiveBFT quick n=7 profile outcome (n=7, f=2, q=5, MaxView=1).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack (n7lite)	5,009	1,454	39	03min 19s	Pass	Safety + W-A*

Snapshot source: make scale-n7lite (artifact: docs/results/scale\_n7lite\_suite\_latest.md, executed on 2026-02-21).

Invariant profile: TypeOKLite + I2-I13 (scalable field-level typing for large constructor domains).

This profile is a bounded quick sanity check, not an unbounded theorem proof.

TABLE V: AdaptiveBFT bounded scale-sanity outcome (n=7, f=2, q=5, MaxView=2).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack (n7)	15,045	4,366	49	04min 50s	Pass	Safety + W-A*

Snapshot source: make scale-n7 (artifact: docs/results/scale\_n7\_attack\_latest.md, executed on 2026-02-21).

Invariant profile: TypeOKLite + I2-I13 (scalable field-level typing for large constructor domains).

This profile is a bounded scale sanity check, not an unbounded theorem proof.

TABLE VI: AdaptiveBFT refinement-transfer TLC outcomes (bounded diagnostics).

Profile	States Generated	Distinct States	Diameter	Time	TLC Result	Checked Properties
baseline (MaxView=2)	269,339	47,796	30	01min 10s	Pass	T0-T4
MaxView=3	269,339	47,796	30	01min 21s	Pass	T0-T4
MaxView=4	269,339	47,796	30	01min 21s	Pass	T0-T4

Commands: make test-transfer, make test-transfer-mv3, make test-transfer-mv4 (executed on 2026-02-21).

T0-T4 definitions are listed in Table II.

Transfer diagnostics are bounded and constraint-scoped; they provide executable evidence for concrete-step to abstract-step/stutter alignment, not a closed unbounded refinement proof.

TABLE VII: AdaptiveBFT wrapper-integrated projection outcomes (baseline and MaxView=4).

Profile	Model	States Generated	Distinct States	Diameter	Time	Checked Projection Properties
baseline	MC_AdaptiveAttack	15,045	4,366	49	11s	W-A*
baseline	MC_LivenessAPS	10,087	3,035	63	13s	W-P*
baseline	MC_LivenessAPSA	45,095	14,737	73	21s	W-J*
MaxView=4	MC_AdaptiveAttack	113,041	32,774	70	43s	W-A*
MaxView=4	MC_LivenessAPS	16,791	5,047	98	39s	W-P*
MaxView=4	MC_LivenessAPSA	104,191	33,269	107	01min 01s	W-J*

Snapshot source: make wrapper-projection (latest artifact: docs/results/wrapper\_projection\_latest.md, executed on 2026-02-21). W-A\*/W-P\*/W-J\* abbreviations are defined in Table II.

These checks are bounded finite-domain diagnostics integrated into the main wrappers.

TABLE VIII: AdaptiveBFT TLC outcomes (bounded model checking, baseline profile).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack	15,045	4,366	49	5s	Pass	Safety only
MC_LivenessAPS	10,087	3,035	63	< 1s	Pass	Safety + L1
MC_LivenessAPSA	45,095	14,737	73	20s	Pass	Safety + L2

Snapshot source: command make matrix (executed on 2026-02-15, after weighted-strike RVS upgrade).

Runtime platform: AMD Ryzen 7 7840H, 16 vCPUs, 15 GiB RAM, Linux (WSL2 Ubuntu), Java -Xmx8G.

Worker policy: MC\_AdaptiveAttack uses 4 workers; liveness models use -workers auto.

These are bounded finite-domain model-checking results, not unbounded theorem proofs.

TABLE IX: AdaptiveBFT deeper-profile TLC outcomes (MaxView=3, bounded model checking).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack	5,936,703	1,196,611	69	7m39s	Pass	Safety only
MC_LivenessAPS	14,762,371	3,405,172	89	59m13s	Pass	Safety + L1
MC_LivenessAPSA	56,392,041	14,807,882	98	3h24m	Pass	Safety + L2

Snapshot source: make test-attack-mv3, make test-aps-mv3, make test-joint-mv3 (executed on 2026-02-15).

Runtime platform: AMD Ryzen 7 7840H, 16 vCPUs, 15 GiB RAM, Linux (WSL2 Ubuntu), Java -Xmx8G.

Worker policy: MC\_AdaptiveAttack uses 4 workers; liveness models use -workers auto.

These are bounded finite-domain model-checking results, not unbounded theorem proofs.

TABLE X: AdaptiveBFT strengthened-profile TLC outcomes (MaxView=4, bounded model checking).

Model	States Generated	Distinct States	Diameter	Time	TLC Result	Claim Scope
MC_AdaptiveAttack	113,041	32,774	70	21s	Pass	Safety only
MC_LivenessAPS	16,791	5,047	97	39s	Pass	Safety + L1
MC_LivenessAPSA	104,191	33,269	106	55s	Pass	Safety + L2

Snapshot source: make test-attack-mv4, make test-aps-mv4, make test-joint-mv4 (executed on 2026-02-15).

Runtime platform: AMD Ryzen 7 7840H, 16 vCPUs, 15 GiB RAM, Linux (WSL2 Ubuntu), Java -Xmx8G.

Worker policy: MC\_AdaptiveAttack uses 4 workers; liveness models use -workers auto.

Strengthened profile parameters keep  $n = 4$ ,  $f = 1$ , and  $q = 3$ , changing only MaxView from 2 to 4.

These are bounded finite-domain model-checking results, not unbounded theorem proofs.